

Machine Learning - Case Study

Introduction

Failed Payments

- Company X is a weekly subscription service. Customers have the option to skip a week's delivery or cancel their account; if they do neither they will be charged each week and sent a box of ingredients and recipes.
- Each week, payment can't be taken for a subset of customers (for example, the credit card on the account is expired). We call these "failed payments".
- Customers with a failed payment are contacted, and have the ability to pay for their order in the following days. However, for logistical reasons a decision must be made within a few hours of the missed payment as to whether to ship the unpaid order
- You've come up with the idea to make a machine learning model that informs the logistics team whether to produce and ship each box that missed payment, based on the likelihood of the customer eventually paying for it.

Data

- You have been provided with two data files:
 - **training_data.csv** for data exploration and building the model
 - **test_data.csv** for model evaluation
- The data files have the following attributes:
 - **customer_id** - a unique id per customer
 - **order_nr** - a unique identifier per order
 - **delivery_date** - the date the order will be delivered (YYYYMMDD)
 - **delivery_week** - the week the order will be delivered
 - **final_payment_status** - whether the customer paid for the order after 10 weeks
 - **time_to_pay** - the time between the failed payment and the customer paying for the order (no units)
 - **payment_method** - the payment method on the customer's account that is charged each week
 - **custom_meal_selection** - whether the customer chose their recipes for the week ("Yes") or received the default selection ("No")
 - **state** - the state of the customer's delivery address
 - **channel** - the marketing channel that the customer signed up through
 - **engagement_score** - represents how engaged the customer is with the brand. Must be > 0.

- ***num_prior_orders_failed*** - the number of previous orders the customer had with a failed payment
- ***num_prior_orders_unpaid*** - the number of previous orders the customer had with a failed payment, that they didn't pay as at the time of the current failed order
- ***avg_recipe_rating*** - the average score the customer gave their past recipes

Part One: Problem Solving

1. The CEO is very interested in your idea to make a model, and would like you to summarise your goal. Write a problem statement: i.e. summarise the challenge you are trying to solve in one sentence/question.

Try to make your problem statement SMART - specific, measurable, achievable, realistic and time-bound.

2. The CEO also wants to know: Do you think shipping orders that fail payments is profitable? Briefly explain your opinion.

Part Two: SQL

Based on the tables below (filled with example data), write SQL queries to get:

1. A list of unique **order_numbers** that failed payment with delivery dates between 2020-02-08 and 2020-02-12 (inclusive) in Australia
2. For each order that failed payment:
 - a. If the order was eventually paid, the number of days between the **first** failed payment and it being paid for by the customer
 - b. If the order was never paid, put **never_paid**
3. **customer_id**, **name**, **phone number** and **order_number** for the most recent order (by delivery date) per customer

Table name = order_payment_history

order_number	time*	status**
1	2020-02-03 10:00	order_created
1	2020-02-03 11:00	payment_charge_attempt
1	2020-02-03 11:01	payment_failed
1	2020-02-04 11:00	payment_charge_attempt
1	2020-02-04 11:01	payment_failed
1	2020-02-05 11:00	payment_charge_attempt
1	2020-02-05 11:01	order_paid
2	2020-02-03 09:00	order_created
2	2020-02-03 11:00	payment_charge_attempt
2	2020-02-03 11:01	order_paid

* yyyy-mm-dd hh:mm

** an order can fail payment multiple times as shown in order 1

'order_paid' status can only appear a maximum of once per order

Table name = orders

customer_id	order_number	delivery_date	country
A	1	2020-02-08 07:00	Australia
B	2	2020-02-12 08:00	Australia
B	3	2020-02-19 08:00	Australia
C	4	2020-02-11 09:00	NZ

Table name = customers

customer_id	name	phone
A	John	0412 345 678
B	Jane	0487 654 321

Part Three: Python

This section should be completed using Python, preferably with Jupyter notebooks.

Key assessment criteria:

- Code structure and readability
- Code reusability (use of functions)
- Thought process explained through comments or text blocks
- Insights extracted

Using the data from **training_data.csv** write Python code to:

1. Perform data cleaning and check data accuracy. These checks will be applied to the test data in question 5 so use functions to make the code reusable.
 - a. Write a function to check data quality
 - b. Write a function to clean the data
2. Write functions to build charts and perform EDA (exploratory data analysis). Provide comments for all charts.
 - a. Produce one chart to examine the target variable and one to examine **time_to_pay**.
 - b. Produce a chart to examine the relationship between **engagement_score** and the target variable (hint: it might be helpful to bucket customers based on the **engagement_score**).
 - c. Produce separate charts to examine the relationship between the target variable and the following features: **num_prior_orders_unpaid**, **custom_meal_selection**, **channel**. The insights might differ between customers with low/high engagement scores - make sure you account for this.
3. Write a function to perform feature engineering:
 - a. Create an **ever_rated_recipe** binary feature based on the **avg_recipe_rating** column
 - b. Create an **is_f** binary feature based on the **channel** (include f1 and f2)
4. Train an ML model to predict whether orders will be paid or unpaid. Use the customer features provided. Note:
 - Treat orders with payment status of Cancelled, Refunded or Other as Paid
 - Include features that you engineered in question 3.
 - Use functions (or other methods of reusable code) to prepare your data for the model.
 - Focus on model evaluation. Evaluate model performance and discuss the results.

Using the data from **test_data.csv** write Python code to:

5. Make predictions using the model built in Q4:
 - a. Check and clean the data (reuse functions from Q1)
 - b. Transform the data (reuse functions from Q4)
 - c. Make predictions
 - d. Evaluate model performance and discuss the results.